



# Real-time Linux

FreedomHEC 2012 Taipei  
Insop Song  
Ericsson



# Contents

- Linux & Real-time
- Linux scheduler
- Latency
- Real-time Linux (PREEMPT\_RT)
- Benchmark test results
- Tools
- Tips
- EDF scheduler
- Summary

# Linux

- Linux is a general purpose OS
  - It's about fair scheduling
  - Main goal is to optimize the average throughput
  - Not deterministic
  - Large latencies



# Real time systems

- Real Time system is ..
  - not “real fast”
  - determinism
- Types of real-time systems are ...
  - Hard real-time
    - Missing a deadline is a total system failure [wiki]
    - Mission critical systems
  - Soft real-time
    - The usefulness of a result degrades after its deadline, thereby degrading the system's quality of service [wiki]
    - Non-mission critical systems



# Linux Scheduler

# Linux scheduler 1

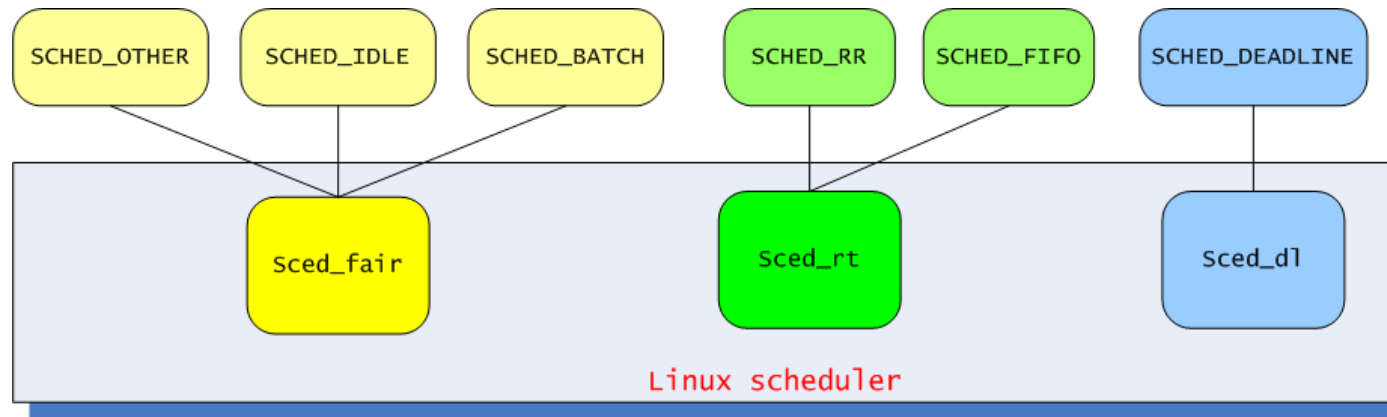
- Linux kernel scheduler
  - Decides which process to run next
  - Allocates processor(s)' time among run-able processes
  - Realizes multi-tasking in a single processor machine
  - Schedules SMP as well
  - Affects how the system behaves, responsiveness

# Linux scheduler 2

- Scheduler framework
  - Kernel supports various scheduling policies by plug-in
  - Scheduling classes contains details of the scheduling policy
    - Decides which task runs next
    - Each operation can be requested by the global scheduler;
    - Allows for creating of the generic scheduler without any knowledge about different scheduler classes

# Linux schedule classes

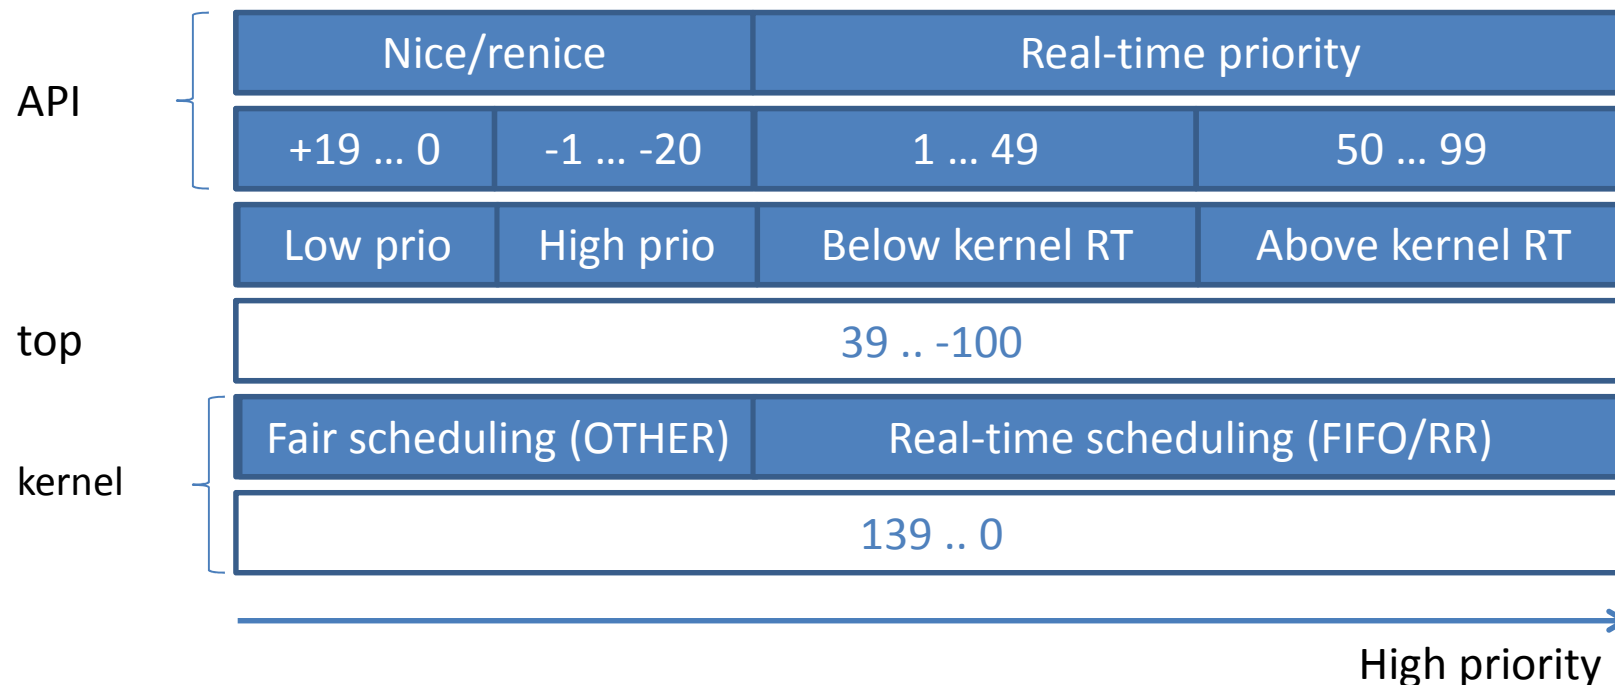
- Supports multiple schedule classes
  - SCHED\_OTHER
    - Default scheduler, non-real-time scheduling
  - SCHED\_FIFO
    - Use FIFO real-time scheduling
  - SCHED\_RR
    - Use round-robin real-time scheduling
  - SCHED\_DEADLINE (not in mainline yet)
    - EDF scheduling with reservation-based scheduler





# Linux scheduler priority

- Two separate priority ranges
  - **Nice value**: -20 ... +19 (19 being the lowest)
  - **Real-time priority**: 0 ... 99 (higher value is higher prio)
  - ***ps* priority**: 0 ... 139 (0: lowest and 139: highest)



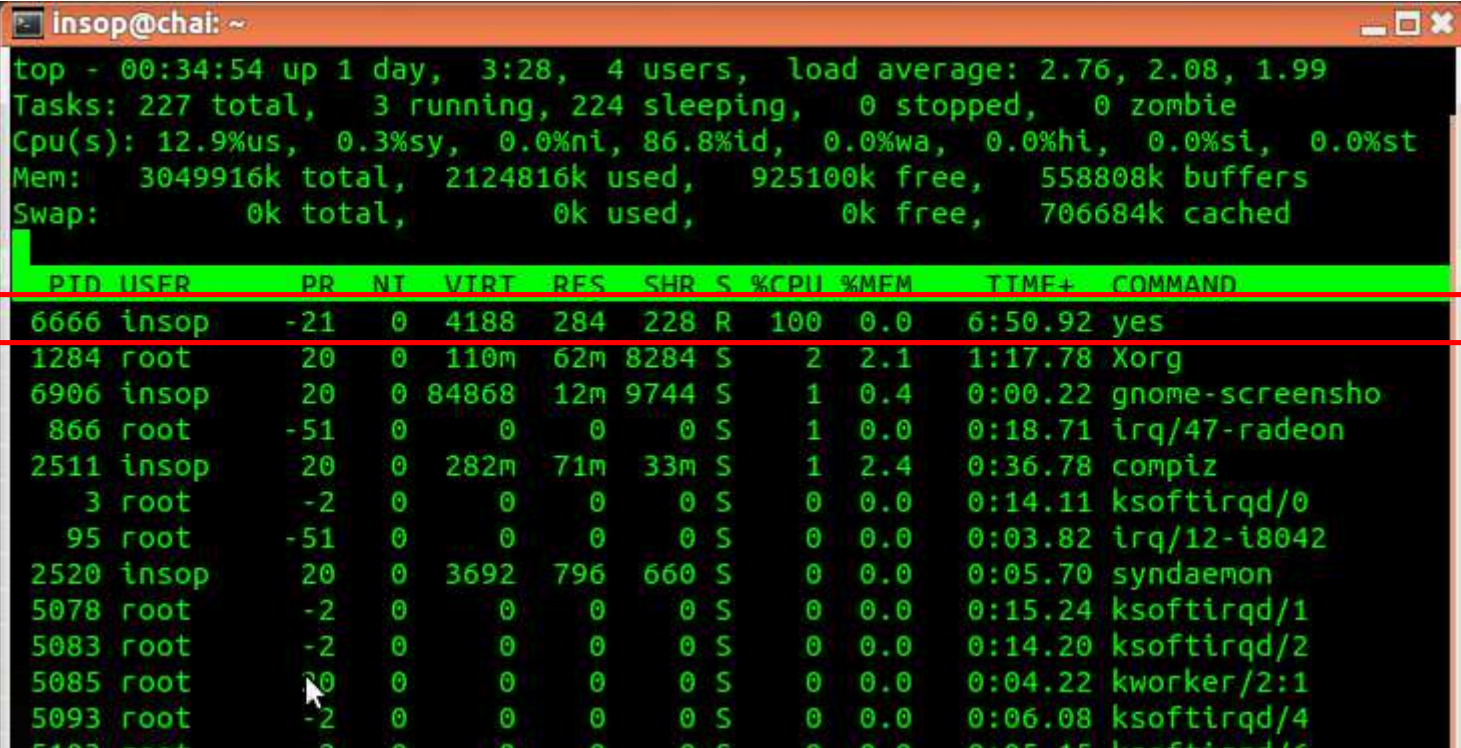
# Linux schedule system calls

- Scheduler related system calls

System call	description
<code>nice()</code>	Sets a process' nice values
<code>sched_setscheduler()</code>	Sets a process' scheduling policy
<code>sched_getscheduler()</code>	Gets a process' scheduling policy
<code>sched_setparam()</code>	Sets a process' real-time policy
<code>sched_getparam()</code>	Gets a process' real-time policy
<code>sched_get_priority_max()</code>	Gets the max real-time priority
<code>sched_get_priority_min()</code>	Gets the min real-time priority
<code>sched_rr_get_interval()</code>	Gets a process' timeslice value
<code>sched_setaffinity()</code>	Sets a process' processor affinity
<code>sched_getaffinity()</code>	Gets a process' processor affinity
<code>sched_yield()</code>	Temporarily yields the processor

# Linux scheduler

- How to config priority
  - `> chrt [option] -p [prio] pid`
    - -r: SCHED\_RR
    - -f: SCHED\_FIFO
  - `sched_setscheduler()`
  - Checking priority, nice, rt-priority
    - `> ps -eo pid,class, pri,nice,rtprio,comm`




```
insop@chal: ~
top - 00:34:54 up 1 day, 3:28, 4 users, load average: 2.76, 2.08, 1.99
Tasks: 227 total, 3 running, 224 sleeping, 0 stopped, 0 zombie
Cpu(s): 12.9%us, 0.3%sy, 0.0%ni, 86.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 3049916k total, 2124816k used, 925100k free, 558808k buffers
Swap: 0k total, 0k used, 0k free, 706684k cached

  PID USER      PR  NI  VIRT  RES  SHR  S  %CPU  %MEM     TIME+  COMMAND
 6666 insop    -21   0   4188   284   228  R   100   0.0    6:50.92  yes
 1284 root       20    0   110m   62m  8284  S     2   2.1    1:17.78  Xorg
 6906 insop     20    0  84868   12m  9744  S     1   0.4     0:00.22  gnome-screensho
  866 root     -51    0      0      0      0  S     1   0.0     0:18.71  irq/47-radeon
 2511 insop     20    0   282m   71m   33m  S     1   2.4     0:36.78  compiz
    3 root      -2    0      0      0      0  S     0   0.0     0:14.11  ksoftirqd/0
   95 root     -51    0      0      0      0  S     0   0.0     0:03.82  irq/12-i8042
 2520 insop     20    0   3692    796   660  S     0   0.0     0:05.70  syn daemon
 5078 root      -2    0      0      0      0  S     0   0.0     0:15.24  ksoftirqd/1
 5083 root      -2    0      0      0      0  S     0   0.0     0:14.20  ksoftirqd/2
 5085 root       0    0      0      0      0  S     0   0.0     0:04.22  kworker/2:1
 5093 root      -2    0      0      0      0  S     0   0.0     0:06.08  ksoftirqd/4
 5103 root      -2    0      0      0      0  S     0   0.0     0:05.15  ksoftirqd/6
```

Latency

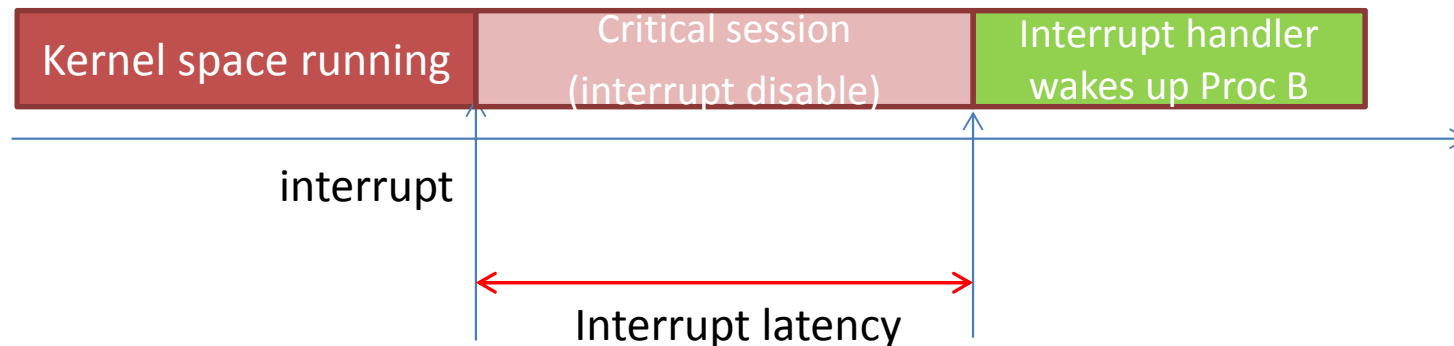
# Latency in Kernel

- Real time means external event should be handled within the bounded time
- Interrupt handler responds to the event and inform user-space process
- Latency
  - Time taken from external interrupt till a user-space process to react to the interrupt



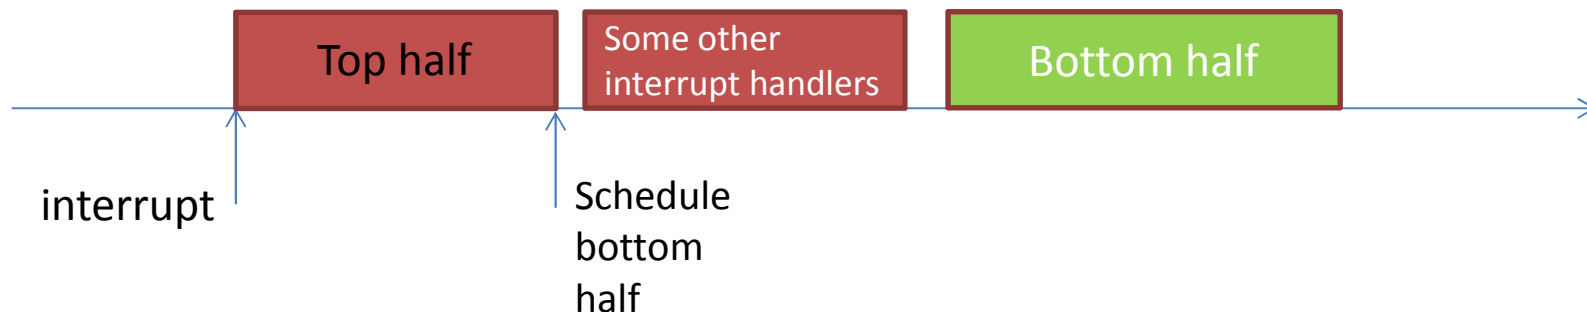
# Source of interrupt latency

- Interrupt latency
  - Time elapsed before the interrupt handle starts
- One of the main cause of the latencies
  - Unbounded critical section (interrupt disabled section)
    - By spinlock and explicit interrupt disable



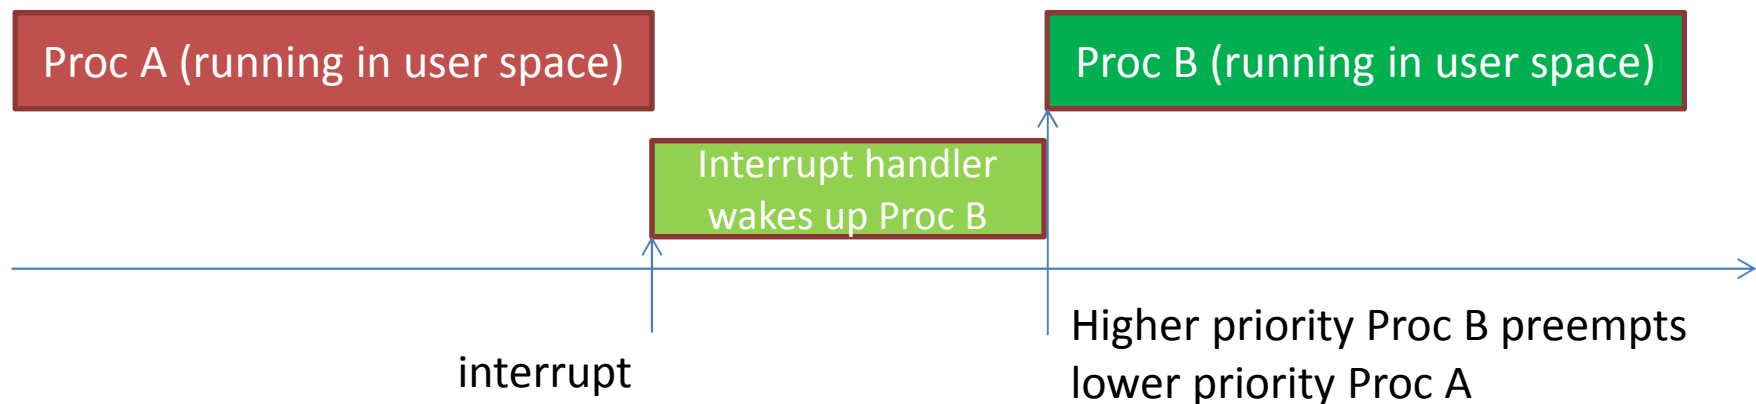
# Interrupt handler

- Interrupt handlers are splitted into two parts
  - A top-half:
    - Process, as quickly as possible, the work during interrupt disabled, such as queue the information for the bottom-half
    - Schedule the bottom-half
  - A bottom-half
    - Process the required tasks from the triggered interrupt



# Kernel preemption 1

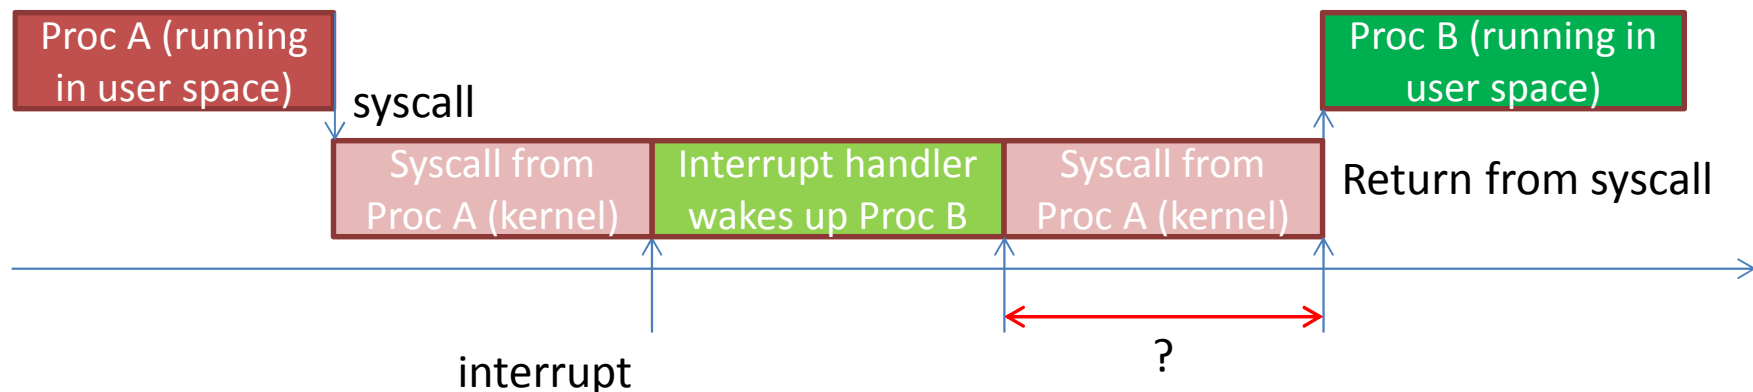
- Kernel can be configured preemptible
- If an interrupt is arrived during Proc A (lower priority) runs
  - Proc A is interrupted
  - Interrupt handler triggers another user space proc (Proc B, higher priority then Proc A)
- When Interrupt handler finished
  - Proc B preempted previously running Proc A





# Kernel preemption 2

- Interrupt triggered during Proc A was invoking syscall, which runs in the kernel space
  - Interrupt handler wakes up Proc B
- However, Proc A running in the kernel space will not be preempted
  - *By default*, the Linux kernel does not do *kernel preemption*
- The time before Proc B runs (scheduler to be called) is unbounded



# Scheduler latency/duration

- Scheduler latency
  - Time elapsed before the scheduler runs
- Scheduler duration
  - Time taken to decide what is the next task to run
  - SCHED\_FIFO & SCHED\_RR use bit map to find out the next task
  - CFS (Complete Fair Scheduler) uses Red-black tree as a sorted queue

Other issues

# Other sources of latencies

- Linux is highly based on virtual memory, as provided by an MMU,
  - so that memory is allocated on demand.  
Whenever an application accesses code or data for the first time, it is loaded on demand, which can create huge delays.
- Many C library services or kernel services are not designed with realtime constraints in mind.

# Real-time Linux

# Real time approaches

- Two major approaches real time Linux
- Approaches
  - Real-time Linux (PREEMPT\_RT patch)
    - Allows preemption, so minimize latencies
    - Execute all activities (including IRQ) in “schedulable/thread” context
    - Many of the RT patch have been merged
  - Linux extention
    - Add extra layer between hardware and the Linux kernel to manage real-time tasks separately
    - (This will not be covered in this presentation)

# Approaches in Real-time Linux

- PREEMPT\_RT
  - A project led by kernel developers including Ingo Molnar, Thomas Gleixner, and Steven Rostedt
    - Large testing efforts at RedHat, IBM, OSADL, Linutronix
  - Goal is to improve real time performance
  - Separate branch, not yet in main-line kernel, but it is getting close
    - Many of small features from PREEMPT\_RT are in main-line kernel over the years

# PREEMPT\_RT

- Background
  - Replace non-preemptible constructs with preemptible ones
- Make OS preemptible as much as possible
  - except preempt\_disable and interrupt disable
- Make Threaded (schedulable) IRQs
  - so that it can be scheduled
- spinlocks converted to mutexes (aka sleeping spinlocks)
  - Not disabling interrupt and allows preemption
  - Works well with thread interrupts



# Approaches in Real-time Linux

- Mainline kernel
  - CONFIG\_PREEMPT option in mainline
  - Allows preemption
  - Mainline 2.6 & 3.0 kernel
    - CONFIG\_PREEMPT\_NONE
      - No forced kernel preemption
    - CONFIG\_PREEMPT\_VOLUNTARY
      - Explicit preemption points in kernel
    - CONFIG\_PREEMPT
      - All kernel code (except critical section) preemptible

# High resolution timers

- Resolution of the system timer (HZ)
  - 100 HZ to 1000 HZ depends on arch and configuration
  - Resolution of 10 ms to 1 ms
  - Adds overheads
- High resolution timers introduced since 2.6.21
  - Allows to use available h/w timers to program interrupts
  - Granularity nano second

# Threaded interrupts

- Handle interrupt by interrupt handler thread
- Interrupt handlers run in normal kernel threads
  - Priorities can be configured
- Main interrupt handler
  - Do minimal work and wake-up the corresponding thread
- Thread interrupts allows to use sleeping spinlocks
- Merged since 2.6.30
  - Conversion is manual in mainline
  - *cf*) in PREEMPT\_RT, all interrupt handlers are switched to threaded interrupt

# Threaded interrupt handler

- Priority must be set
  - Interrupt handler threads
  - Softirq threads
  - Other kernel threads
  - Real time application processes/threads

# PREEMPT\_RT

- Almost all kernel space is now preemptible
  - An interrupt can occur at any time
  - The woken up processes by interrupt can start immediately
- Treaded IRQs
  - Kernel thread per ISR

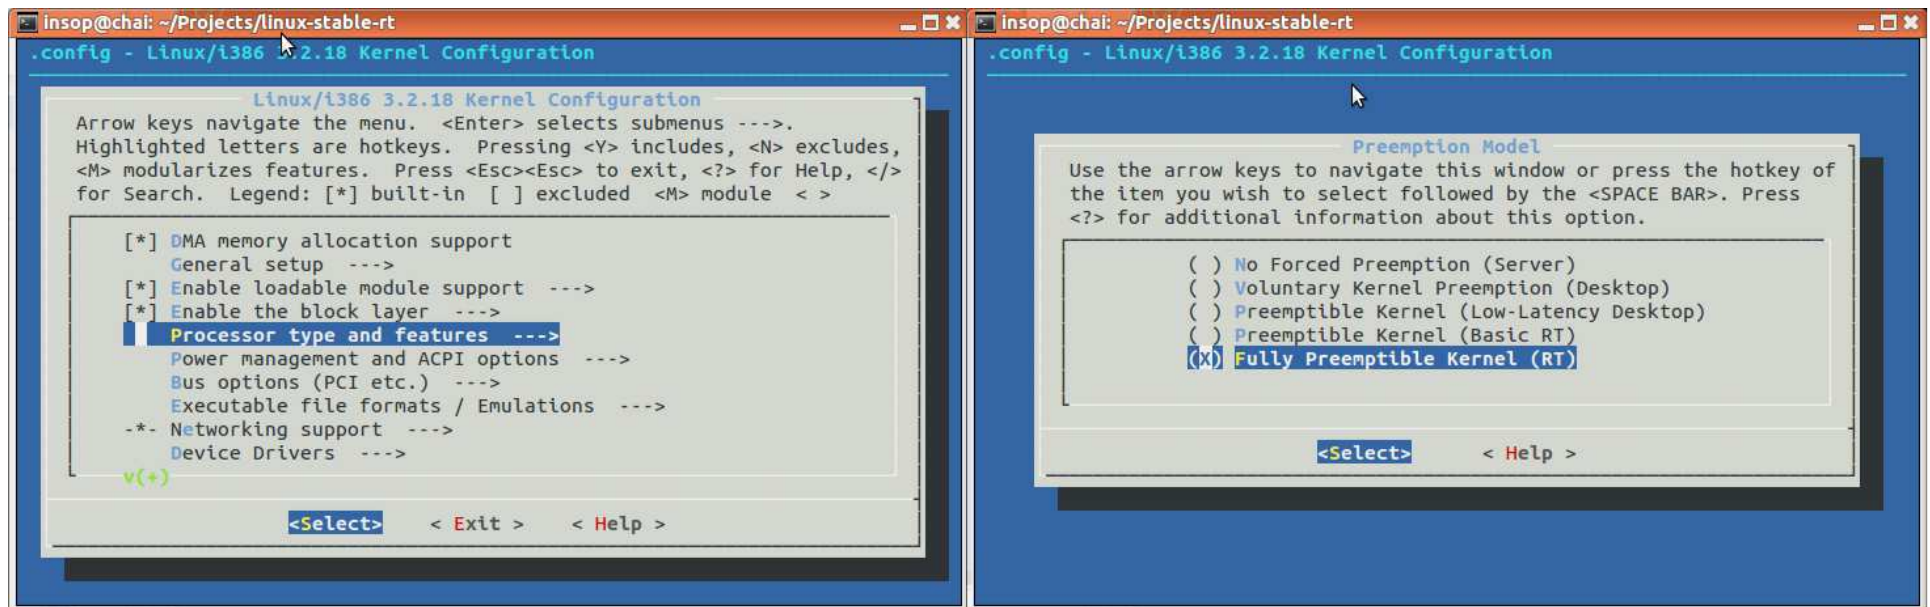
# Configuring PREEMPT\_RT

# PREEMPT\_RT config

- Where to get:
  - <http://www.kernel.org/pub/linux/kernel/projects/rt/>
- Pre 3.0
  - Use `-rt` patch, i.e. download mainline & apply patch
- Post 3.0
  - Patch as well as git available
    - [git clone git://git.kernel.org/pub/scm/linux/kernel/git/rt/linux-stable-rt.git](http://git.kernel.org/pub/scm/linux/kernel/git/rt/linux-stable-rt.git)
- Yocto project supports PREEMPT\_RT as well
  - Reference images
    - `core-image-rt`
    - `core-image-rt-sdk`
  - Detail information:
    - <http://www.yoctoproject.org/docs/current/poky-ref-manual/poky-ref-manual.html>

# PREEMPT\_RT config

- Processor type and features
  - Preemption model
    - Fully Preemptible Kernel (RT)
  - High-resolution timer support
  - disabled all Power Management Options like ACPI or APM





# RT hello world

```
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <sched.h>
#include <sys/mman.h>
#include <string.h>

#define MY_PRIORITY (49) /* we use 49 as the PRREMPT_RT use 50
                           as the priority of kernel tasklets
                           and interrupt handler by default */

#define MAX_SAFE_STACK (8*1024) /* The maximum stack size
which is
                                guranteed safe to access without
                                faulting */

#define NSEC_PER_SEC (1000000000) /* The noof nsecs per sec. */
void stack_prefault(void) {
    unsigned char dummy[MAX_SAFE_STACK];
    memset(dummy, 0, MAX_SAFE_STACK);
    return;
}
...
```

[https://rt.wiki.kernel.org/index.php/RT\\_PREEMPT\\_HOWTO](https://rt.wiki.kernel.org/index.php/RT_PREEMPT_HOWTO)

```
int main(int argc, char* argv[])
{
    struct timespec t;
    struct sched_param param;
    int interval = 50000; /* 50us */
    param.sched_priority = MY_PRIORITY; /* Declare ourself as a real time task */
    if(sched_setscheduler(0, SCHED_FIFO, &param) == -1) {
        perror("sched_setscheduler failed");
        exit(-1);
    }
    /* Lock memory */
    if(mlockall(MCL_CURRENT|MCL_FUTURE) == -1) {
        perror("mlockall failed");
        exit(-2);
    }
    stack_prefault(); /* Pre-fault our stack */
    clock_gettime(CLOCK_MONOTONIC, &t);
    /* start after one second */
    t.tv_sec++;
    while(1) {
        /* wait until next shot */
        clock_nanosleep(CLOCK_MONOTONIC, TIMER_ABSTIME, &t, NULL);
        /* do the RT stuff */
        t.tv_nsec += interval; /* calculate next shot */
        while (t.tv_nsec >= NSEC_PER_SEC) {
            t.tv_nsec -= NSEC_PER_SEC;
            t.tv_sec++;
        }
    }
}
```

# PREEMPT\_RT benchmark results

# Benchmark result of vanilla kernel

- Cyclictest
  - measuring accuracy of sleep and wake operations of highly prioritized realtime threads
  - <https://rt.wiki.kernel.org/index.php/Cyclictest>
  - `git clone git://git.kernel.org/pub/scm/linux/kernel/git/clkwlms/rt-tests.git`
  - `sudo ./cyclictest -a -t -n -p99`

Vanilla kernel: 3.2.0-24 on 8 core PC, CONFIG\_PREEMPT\_VOLUNTARY

```
insop@chai:~/Projects/rt-tests$ uname -a
Linux chai 3.2.0-24-generic-pae #39-Ubuntu SMP Mon May 21 18:54:21 UTC 2012 i686 i686 i386 GNU/Linux
insop@chai:~/Projects/rt-tests$ sudo ./cyclictest -a -t -n -p99
# /dev/cpu_dma_latency set to 0us
policy: fifo: loadavg: 0.54 0.69 0.67 6/417 3256

T: 0 ( 2772) P:99 I:1000 C:1008249 Min:      4 Act:   18 Avg:   11 Max:    701
T: 1 ( 2773) P:99 I:1500 C: 672166 Min:      4 Act:   35 Avg:   11 Max:    491
T: 2 ( 2774) P:99 I:2000 C: 504124 Min:      4 Act:    9 Avg:   11 Max:    363
T: 3 ( 2775) P:99 I:2500 C: 403299 Min:      4 Act:   14 Avg:   11 Max:   2013
T: 4 ( 2776) P:99 I:3000 C: 336082 Min:      4 Act:   14 Avg:   14 Max:    804
T: 5 ( 2777) P:99 I:3500 C: 288071 Min:      3 Act:   13 Avg:    9 Max:    190
T: 6 ( 2778) P:99 I:4000 C: 252062 Min:      3 Act:    9 Avg:    9 Max:    343
T: 7 ( 2779) P:99 I:4500 C: 224055 Min:      3 Act:   13 Avg:   10 Max:    224
```

Worst case latency: hundreds of usec

# Benchmark result of PREEMPT\_RT

- Cyclictest
  - measuring accuracy of sleep and wake operations of highly prioritized realtime threads
  - <https://rt.wiki.kernel.org/index.php/Cyclictest>
  - git clone git://git.kernel.org/pub/scm/linux/kernel/git/clrkwlms/rt-tests.git
  - *sudo ./cyclictest -a -t -n -p99*

RT kernel: 3.2.18-rt on 8 core PC

```
insop@chai:~/Projects/rt-tests$ uname -a
Linux chai 3.2.18-rt #1 SMP PREEMPT RT Fri May 25 23:04:58 PDT 2012 i686 i686 i386 GNU/Linux
insop@chai:~/Projects/rt-tests$ sudo ./cyclictest -a -t -n -p99
[sudo] password for insop:
# /dev/cpu_dma_latency set to 0us
policy: fifo: loadavg: 0.69 0.79 0.77 1/434 3037

T: 0 ( 2995) P:99 I:1000 C:1030921 Min:      5 Act:      7 Avg:      12 Max:      32
T: 1 ( 2996) P:99 I:1500 C: 687280 Min:      5 Act:      7 Avg:      13 Max:      53
T: 2 ( 2997) P:99 I:2000 C: 515455 Min:      4 Act:      6 Avg:      13 Max:      34
T: 3 ( 2998) P:99 I:2500 C: 412364 Min:      5 Act:      7 Avg:      13 Max:      34
T: 4 ( 2999) P:99 I:3000 C: 343637 Min:      6 Act:     11 Avg:      16 Max:      31
T: 5 ( 3000) P:99 I:3500 C: 294546 Min:      3 Act:      4 Avg:      11 Max:     338
T: 6 ( 3001) P:99 I:4000 C: 257727 Min:      4 Act:      5 Avg:      11 Max:      24
T: 7 ( 3002) P:99 I:4500 C: 229091 Min:      3 Act:      5 Avg:      11 Max:      29
```


Worst case latency: tens of usec

# Benchmark result of vanilla kernel 2

- Cyclicttest running 1 hour

– *./cyclicttest -p 80 -n -D 1h -v | gzip > /tmp/cyclic-rt.gz*

Vanilla kernel: 3.2.0-24, CONFIG\_PREEMPT\_VOLUNTARY

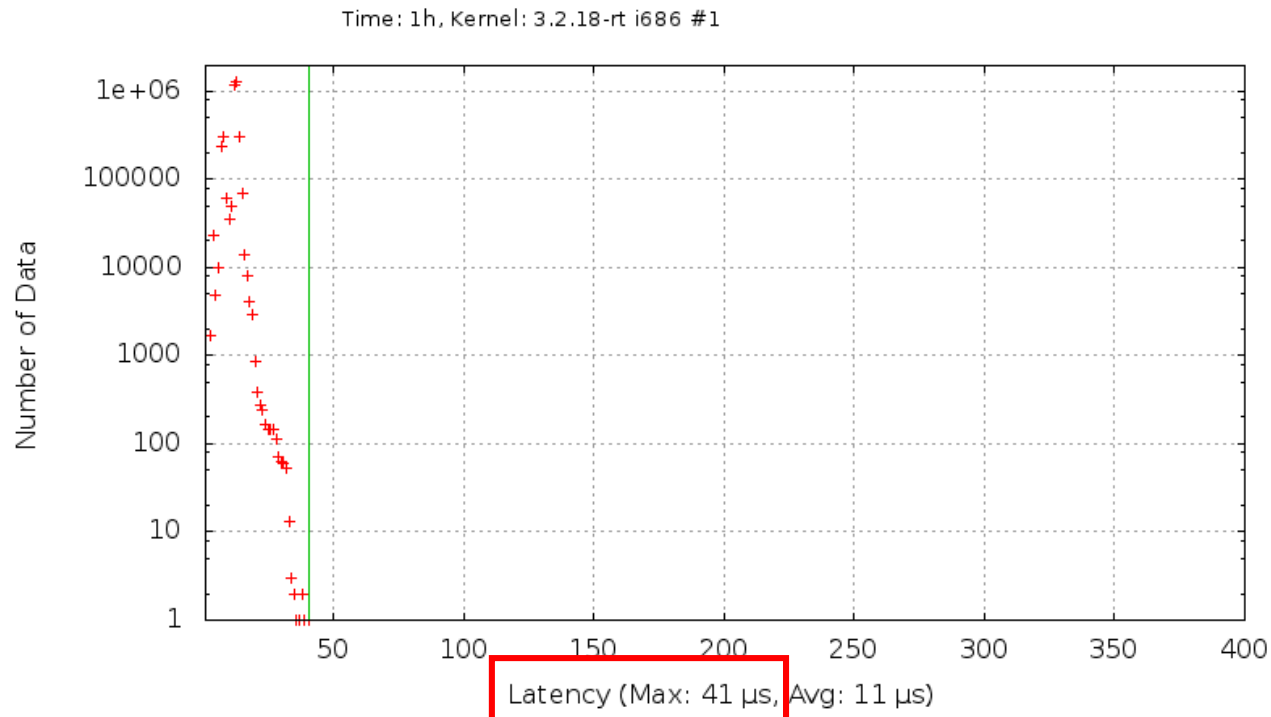


# Benchmark result of PREEMPT\_RT 2

- Cyclicttest running 1 hour

– *./cyclicttest -p 80 -n -D 1h -v | gzip > /tmp/cyclic-rt.gz*

RT kernel: 3.2.18-rt




# Benchmark result of vanilla kernel 3

- Cyclictest running 2 hours

— *./cyclictest -p 80 -n -D 2h -v | gzip > /tmp/cyclic-rt.gz*

Vanilla kernel: 3.2.0-24, CONFIG\_PREEMPT\_VOLUNTARY




# Benchmark result of PREEMPT\_RT 3

- Cyclictest running 2 hours

— *./cyclictest -p 80 -n -D 2h -v | gzip > /tmp/cyclic-rt.gz*


RT kernel: 3.2.18-rt





# Benchmark result from OSADL

- QA farm running at OSADL.ORG
  - QA FARM: <https://www.osadl.org/QA-Farm-Realtime.qa-farm-about.0.html>
  - Comparison between 2.6.32 and 2.6.33-rt
  - <https://www.osadl.org/Compare-systems.qa-farm-compare-latency.0.html>
  - Worst case latency in **non-rt is over 4,000 usec**, in **rt around 100 usec**
  - The following is worst case latency distribution comparisons



# Few selected deployment examples

- US Navy Zumwalt
  - Total Ship Computing Environment (TSCE)
- US Navy VTOL drone (MQ8B Fire Scout)
  - Ground control by uncle Linus who's replacing uncle Gates suffering from various security issues
- Laser for cutting welding
  - TRUMPF Laser Division
- Robots
  - Lot of them [[here](#)]



# PREEMPT\_RT Roadmap

- Complete full PREEMPT\_RT patch merge to mainline
  - *“And yeah, I still think the hard-RT people are mostly crazy”* by Linus 2010
    - (read) getting close, but might take little more time ...
- However,..
  - *“The RT people have actually been pretty good at slipping their stuff in, in small increments, and always with good reasons for why they aren't crazy”* by Linus 2010
    - (read) small and well defined piecewise stuff can be merged ...

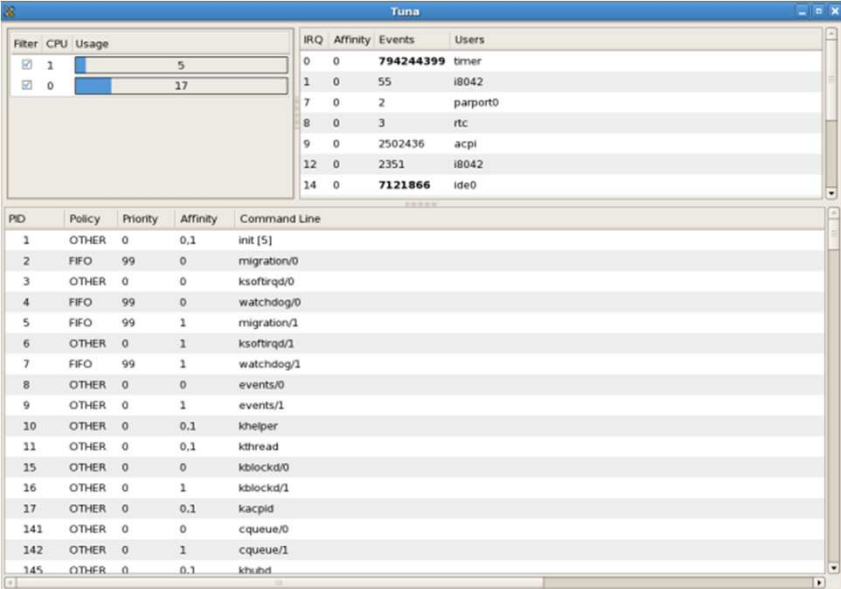
From <http://www.redhat.com/archives/utrace-devel/2010-January/msg00319.html>

# Performance tuning tools

- Ftrace
  - Developed by Steven Rostedt at Redhat
  - Well documented in [Documentation/ftrace.txt](#)
  - Run time trace capture of longest latency paths, kernel and user space
  - Peak detector
  - Detailed kernel profiles
    - Can be used to trace any kernel function

# Performance tuning tools

- TUNA
  - Redhat MRG Realtime
  - Dynamic parameter tunings including process affinity, scheduling policy, device IRQ priority
- Linux performance tools
  - Oprofile: system level statistical profileing



The screenshot shows the Tuna application window. It has two main panes. The top pane displays a table of process affinity and scheduling parameters. The bottom pane displays a table of process affinity and scheduling parameters.

Filter	CPU	Usage
<input checked="" type="checkbox"/>	1	5
<input checked="" type="checkbox"/>	0	17

IRQ	Affinity	Events	Users
0	0	794244399	timer
1	0	55	i8042
7	0	2	parport0
8	0	3	rtc
9	0	2502436	acpi
12	0	2351	i8042
14	0	7121866	ide0

PID	Policy	Priority	Affinity	Command Line
1	OTHER	0	0.1	init [5]
2	FIFO	99	0	migration/0
3	OTHER	0	0	ksoftirqd/0
4	FIFO	99	0	watchdog/0
5	FIFO	99	1	migration/1
6	OTHER	0	1	ksoftirqd/1
7	FIFO	99	1	watchdog/1
8	OTHER	0	0	events/0
9	OTHER	0	1	events/1
10	OTHER	0	0.1	khelper
11	OTHER	0	0.1	kthread
15	OTHER	0	0	kblockd/0
16	OTHER	0	1	kblockd/1
17	OTHER	0	0.1	kacpid
141	OTHER	0	0	cqueue/0
142	OTHER	0	1	cqueue/1
145	OTHER	0	0.1	kthubd

# Practical tips

- Tune the system
  - Tweak priority
    - `chrt`, `sched_setscheduler()`
    - `SCHED_FIFO` and `SCHED_RR` for real-time scheduling
    - Apply RM (Rate monotonic) scheduling scheme
      - Assign higher priorities to more frequently running process
    - Avoid `SCHED_FIFO` priority 99
  - Find hotspot
    - `oprofile`, `gprofile`
  - Find the long latency
    - Measure and estimate worst case response time
    - `Ftrace`
- Processor affinity
  - Bind threads to specific processors
  - Use `cpu affinity` field in `/proc/irqs/<n>/smp_affinity` to bind interrupts to specific processors

# Response time and throughput

- Overhead for real-time preemption
  - Mutex instead of spin lock
  - Priority inheritance
  - Threaded interrupt handler
- Due to overhead of real-time preemption
  - Throughput is reduced
- Due to the flexibility of preemption
  - Much better worst case latency


# Deadline scheduler



# Deadline scheduler


- Optimal dynamic-priority scheduling algorithm
  - A task with a shorter deadline has a higher priority
  - Executes a job with the earliest absolute deadline
- Be able to schedule periodic and aperiodic tasks
- Real-time system is schedulable under EDF if and only if
  - $\sum C_i/T_i \leq 1$ 
    - where  $C_i$ : worst-case execution time
    - where  $T_i$ : period of task  $i$
  - Full processor utilization

[5]



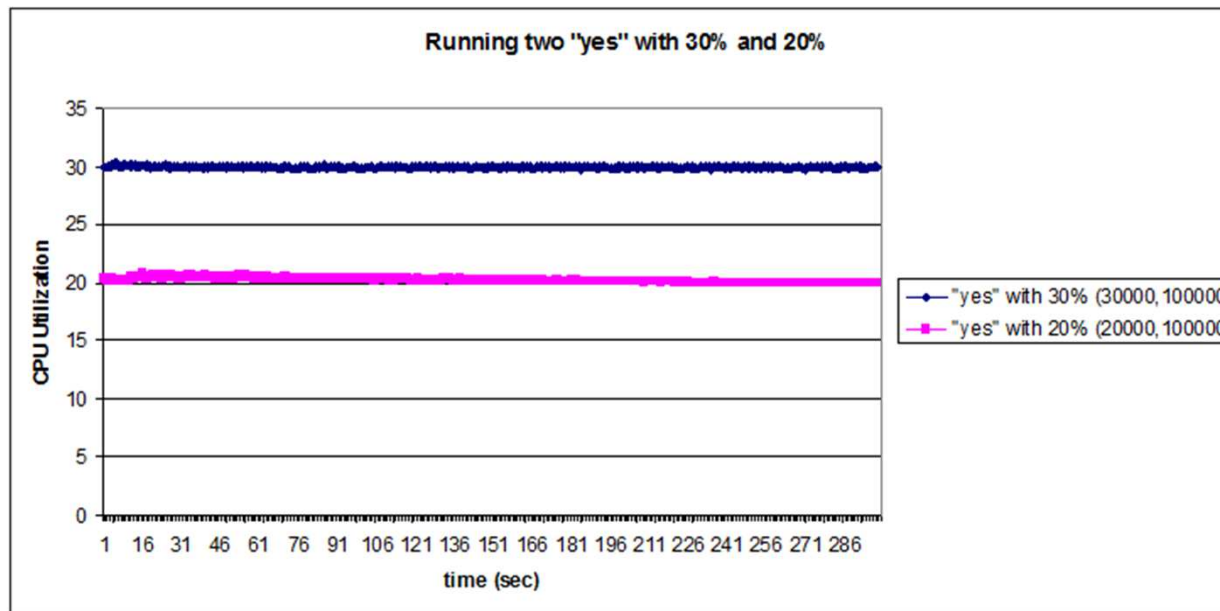
# Deadline scheduler

- Enables a task to be reserved for fractions of the full CPU resource
  - Assigns percentage of CPU to a task instead of assigning priority
  - Ensures each task can only use up to the assigned percentage of the processor even if it demands more
- Guarantees temporal isolation
  - Tasks are temporally isolated, i.e. misbehaving tasks will not impact normal task



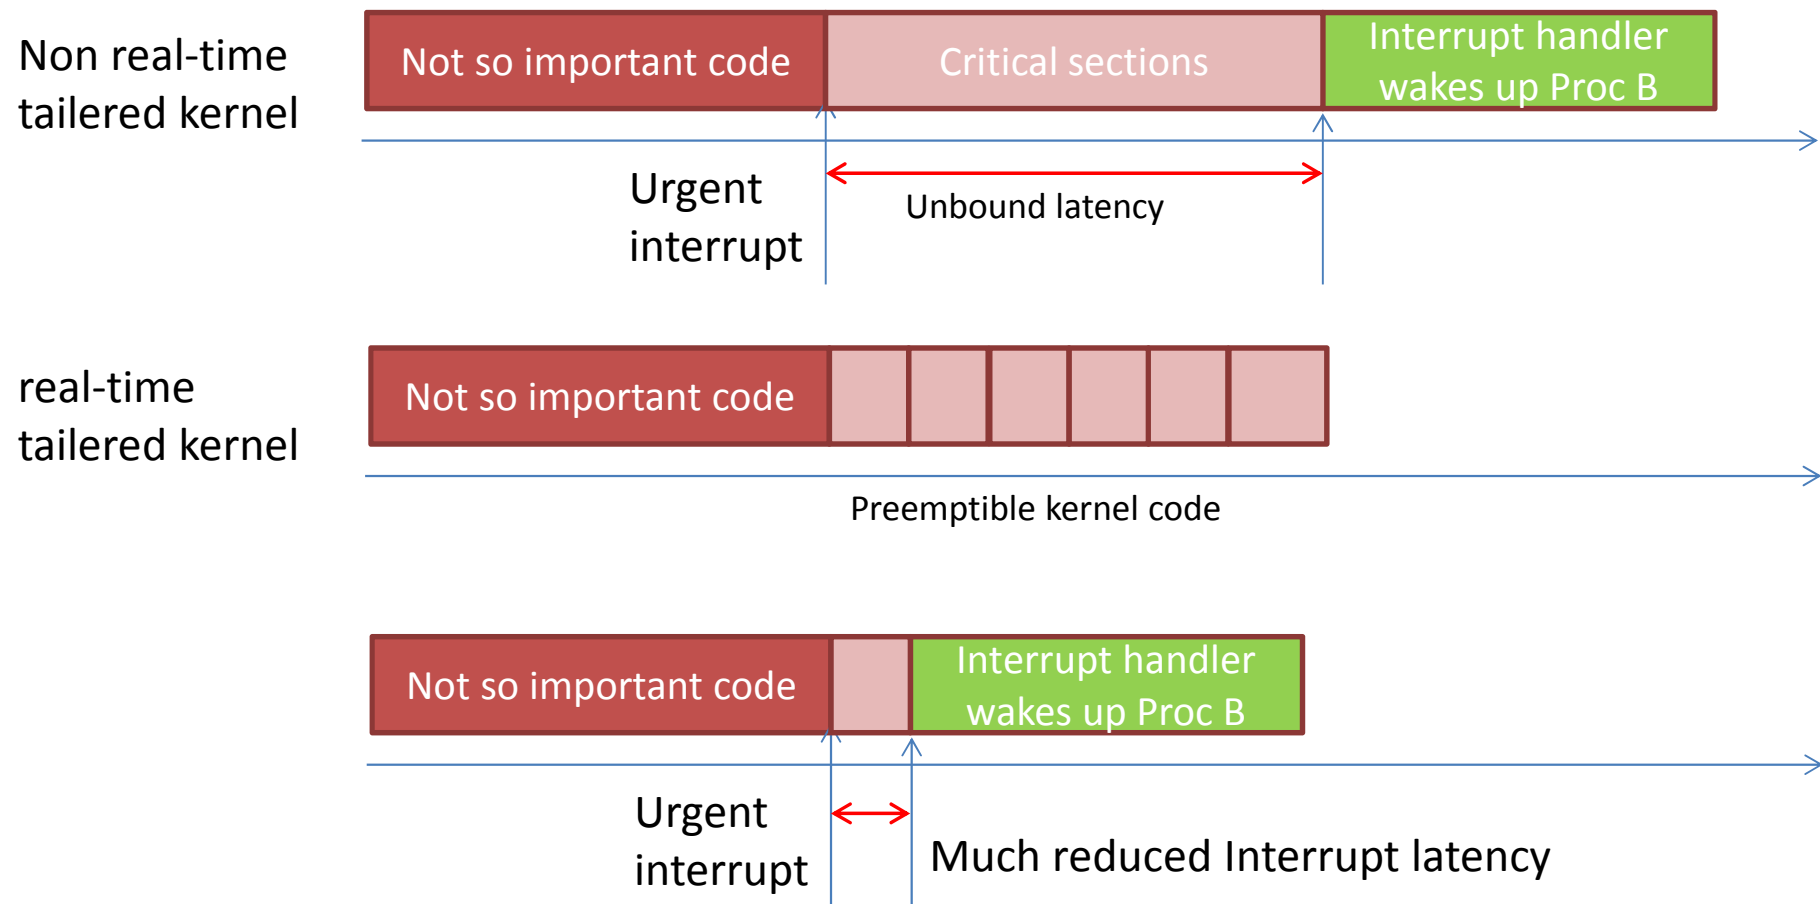
# Deadline/reservation test result

- Running two “yes” with 30% and 20%
  - Yes1: 30%, Q:30000, P:100000
  - Yes2: 20%, Q:20000, P:100000



# Summary

- Most important aspects of Real-time
  - Controlling latency by allowing kernel to be preemptible everywhere..



# Summary

- Mainline kernel
  - Preemptible
    - CONFIG\_PREEMPT
  - Real-time scheduler class
    - SCHED\_FIFO, SCHED\_RR with proper priority
- PREEMPT\_RT patch
  - Virtually most of the kernel code is preemptible
  - Many of features are already merged
    - Many kernel developers from various companies have long been working
    - It is anticipated that more code will be merged to mainline
  - Better worst time latency
  - *“So I can work with crazy people, that's not the problem. They just need to sell their crazy stuff to me using non-crazy arguments, and in small and well-defined pieces.”* Linus

# References

1. [Linux-rt: Turning a General Purpose OS into a Real-Time OS](#), Peter Zijlstra
2. [Inside The RT Patch](#), Steven Rostedt
3. [Real-Time Preemption Patch-Set](#), Manas Saksena
4. [Realtime in embedded Linux systems](#), Michael Opdenacker et al. (free-electrons.com, largely based on this material)
5. [Reservation-Based Scheduler for Network Equipments](#), Insop Song
6. [Linux Kernel Development 3rd Edition](#), Robert Love
7. [Professional Linux Kernel Architecture](#), Wolfgang Mauerer
8. [Introduction to realtime linux](#), Bryan Che, SCaLE, Feb 2009
9. Using Real-time Linux, Klaas van Gend, ELC 2008
10. Linux Realtime-Fähigkeiten, Stefan Agner, 2009
11. [Real-time Track, Linux Plumbers Conference](#), Boston, 2010

# Web links

1. Real-Time Linux Wiki,

<https://rt.wiki.kernel.org/>

2. Mainline RT-preempt patchset,

[http://elinux.org/Mainline RT-preempt patchset](http://elinux.org/Mainline_RT-preempt_patchset)

# Disclaimer

- This work represents the view of the author and does not necessarily represent the view of Ericsson



- Thank you for your attention.
- Question?

Backup

# Real-time Linux

- How to make linux real-time
  - Minimize interrupt disable time
  - Interrupt handling with schedulable threads
  - Almost fully preemptible kernel
    - except short critical sections
  - Synchronization using mutexs instead of spin locks
    - Allows preemptible
  - Priority inheritance support for mutex
  - Priority queue supported mutex
  - High resolution timers

# Kernel config options

- CONFIG\_PREEMPT
- Kernel preemption
- Better Posix real-time API
- Priority inheritance support for Mutexes
- High-resolution timers
- Threaded interrupts
- Spinlock annotations

# Priority inversion

- Priority inversion

# Priority inheritance

- Priority inheritance